

Analyzing Elo Ratings For NCAA Men's Division 1 Hockey

Alex Tidd

Advisors: Matt Higham and Robin Lock

2025-05-02

Table of contents

1 Things to note	1
2 Abstract.....	1
3 Introduction	2
4 Data	3
5 Elo Model	3
6 Results	5
7 Conclusion	14
8 Code Appendix	14
9 Works Cited	23
9.1	23

1 Things to note

- APA Citation
- Render as docx then pdf
- 1.5 spacing or double space
- add line numbers. On rendered file.

2 Abstract

There are 64 NCAA Division 1 Men's Hockey teams. In NCAA Division 1 Men's Hockey, the U.S. College Hockey Online (USCHO) provides the official rankings on the 64 total teams using a system that relies on expert votes. However, this method is not perfect as there is no formal quantitative analytics involved in the voting. To improve on this, a chess ratings system, called Elo, is commonly used in many sports to quantitatively rate players, weighing strength of opponent and recency of match. In this project, we modify the Elo system for NCAA Division 1 Men's Hockey by adding in weights for game goal differential (so that games in which the score differential was large potentially result in a larger bump

in Elo for the winning team) and home-ice advantage (so that the team playing on home-ice has an adjusted probability of winning the game).

3 Introduction

- Goal
- Outline of write up
- Background of Elo

NCAA Men's Division 1 Ice Hockey has 64 teams. This inherently makes ranking teams extremely hard. It is impossible to rank teams in the same manner as pro leagues like the NHL and junior leagues like the 3 CHL leagues, where teams are ranked based off of a record-point system.

To rank teams, the current method is an "expert vote" system done by US College Hockey Online (USCHO). This method takes votes by "experts" and ranks teams based on how many votes each team receives. In general, these rankings are pretty accurate, however it lacks any true quantitative analysis, instead relying on the opinions of "experts".

This project aims to solve the issue presented by using a ratings system used in chess to accurately rate and rank players, called Elo, and use in for collegiate hockey. In chess, Elo ratings calculate an "expected outcome" using 2 players' ratings. This expected outcome is essentially a "win probability" for each player. After the expected outcome is calculated, it is compared to the actual outcome. Depending on the outcome, a players new rating will either be increased or decreased. The benefit to using Elo, is that the rating system takes into consideration your strength, the opponents strength and the recency of the match. Big wins/losses garner big adjustments in rating, whereas an expected win/loss won't inflate/tarnish a rating. The equations are shown below:

Expected outcome:

$$E_a = \frac{1}{1 + 10^{\frac{R_b - R_a}{400}}}$$

Rating update:

$$R'_a = R_a + k(\text{outcome} - E_a)$$

In this case, k is an "update factor" that scales how many points are added or subtracted from a team's rating after the event of a win or loss.

The goal for this project is to use this base Elo ratings system, incorporate factors such as home ice advantage and goal differential, as well as optimize k, to make a ratings model that can accurately predict outcome and rank teams accordingly. This paper will take you through the steps of finding data, creating a ratings function, and optimizing our constants.

4 Data

There were two main data sets used for this project. The full 2023-2024 NCAA Men's Division 1 Ice Hockey schedule, which after wrangling, included 1166 games with variables: date, game_type, away_team, away_score, home_team, home_score, overtime, neutral_site, score_diff, and outcome. The score_diff and outcome variables are in reference to the home team, with outcome being either 1: win, 0.5: tie, or 0: loss. The next data frame was the entire 2024-2025 NCAA Men's Division 1 Ice Hockey schedule. The schedule contains 1153 games which had the same variables and setup as the 2023-2024 season. Both data frames were scraped from the website: College Hockey News.

A third data frame for initial rankings was created from scratch using knowledge from the final 2022-2023 season. Initial rankings were created by ranking teams from 2000 to 1300. The top 8 teams from the 2022-2023 season were assigned the rating of 2000, every 8 teams, the assigned rating would decrease by 100. This initial ratings file was used as the initial ratings file to run with the 2023-2024 season to get accurate ratings for the 2024-2025 season.

To get these "accurate" rating for the 2024-2025 schedule. The 2023-2024 schedule was run through the Elo function with $k = 100$, a value used based off a general exploration which gave relatively accurate rankings compared to the final 2023-2024 season. After, the ratings were scaled using the equation below:

$$\text{rating}^{\wedge} = (\text{rating} \cdot 0.7) + 450$$

This adjustment of rankings is used to adjust for player turnover, new coaching hires, and off season improvement. The formula used is based off of the formula used by Fivethirtyeight in their NHL Elo ratings. These final adjusted ratings are what was used for the initial rating in the 2024-2025 season.

5 Elo Model

To add goal differential into the Elo system, it was determined that it should be placed in the "rating update" portion of the Elo system, due to goal differential being a "post-game" statistic. On the other hand, home ice advantage affects the predicted probability of a home team winning, therefore, it was added to the "expected outcome calculation" portion of the Elo system. Home ice advantage was also incorporated as a simple point "boost" for the home team instead of a multiplier, based off of FiveThirtyEight's NHL Elo model.

In order to obtain the "best" values for each parameter, a grid search method was used to optimize each of the three parameters to lower the mean absolute residual of the season. In this case the "absolute residual" is calculated by the absolute value of the difference of the game outcome and the expected outcome of the home team. Essentially we want the smallest difference between actual outcome and expected outcome. The smaller the difference, the more accurate our function is predicting actual outcomes of games.

103 The un-optimized Elo function is as follows:

104
$$E_{home} = \frac{1}{1 + 10^{\frac{R_{away} - (R_{home} + home_{ice})}{400}}}$$

105
$$E_{away} = \frac{1}{1 + 10^{\frac{(R_{home} + home_{ice}) - R_{away}}{400}}}$$

106
$$R_{home} = R_{home} + k(d(\text{scoreDiff}) + (\text{outcome} - E_{home}))$$

107
$$R_{away} = R_{away} + k(d(-\text{scoreDiff}) + (\text{outcome} - E_{away}))$$

108 Using this model a general exploration to see what possible values of k, d, and home_ice
109 could render a lower mean absolute residual. Based off the exploration, k is found at
110 roughly 100, d at roughly 50, and home_ice at roughly 50. It should be noted that k was
111 explored first with d = 0, and home_ice = 0. Next was home_ice, which was found at varying
112 weights of k and d = 0. d was then explored with k = 100, and home_ice = 0. the only reason
113 k is held in all three cases, is that if k = 0, the function would not be able to update rating as
114 the update portion would always be the initial rating + 0.

115 These explorations provided essential knowledge for the grid search as initial optimization
116 could be found using smaller ranges of values, significantly decreasing the run time of
117 each optimization. This background knowledge cut run time of one optimization dawn from
118 16.3 days to roughly 30 minutes as each optimization grid search used 1000 combinations
119 of the three variables instead of 4.5 million.

120 Using this method k was found to be optimized at 37.22222, home_ice = 53.33333, and d =
121 40.55556, this provided a mean absolute residual of 0.3515844. This looked very promising
122 at first glance. However, upon deeper inspection, two major concerns were found. First,
123 the bottom four teams had end of season ratings in the negatives. Second, there were
124 many residuals of 1 and 0. This raised the most concern as expected outcome is bounded
125 by $0 < \text{expected outcome} < 1$, meaning we should never see residuals of exactly 1 or 0.
126 Upon deeper inspection, it was found that there were many games were expected
127 outcome was to the effect of 0.9 repeating or 1e-20. Which whilst technically within our
128 bounds, is highly unrealistic. A good model that predicts win probability should never give
129 a pregame win probability of nearly 100% or 0%. Something needed to be changed with the
130 model. Goal differential was looked at first.

131 Upon inspecting and researching different ways to incorporate goal differential in an Elo
132 function, FiveThirtyEight was looked at again. In their NHL model, goal differential was
133 incorporated as shown below:

134
$$(0.6686 \cdot \ln(\text{scoreDiff})) + 0.8048$$

135 The new update function now looks as follows:

```
136 $$R^`_{home} = R_{home} + k(((0.6686 \cdot \ln(\text{scoreDiff}))+0.8048)(\text{outcome} -  
137 E_{home}))$$
```

```
138 $$R^`_{away} = R_{away} + k(((0.6686 \cdot \ln(\text{-scoreDiff}))+0.8048)(\text{outcome} -  
139 E_{away}))$$
```

140 The addition of 0.8048 is put in as the factor for score_diff = 1 as the natural logarithm of 1
141 is 0, which would then render the entire update portion of the function useless, running
142 into the same issue if k = 0.

143 Now that goal factor is optimized, all that was left was to rerun the grid search to optimize k
144 and home_ice. Using the grid search method, k is optimized at 88 and home_ice optimized
145 at 40, which gives a mean absolute residual of 0.3876531. Whilst this may not have a mean
146 absolute residual lower than the first model, it maintains that all teams keep ratings above
147 zero and keep predicted outcomes in a more realistic zone.

148 6 Results

149 Results from this optimization show that k = 88, home_ice = 40, and the calculation to
150 utilize score differential = $(0.6686 \cdot \ln(\text{scoreDiff})) + 0.8048$. The plots below were used to
151 further imply this.

152 Grid search for k and home_ice:

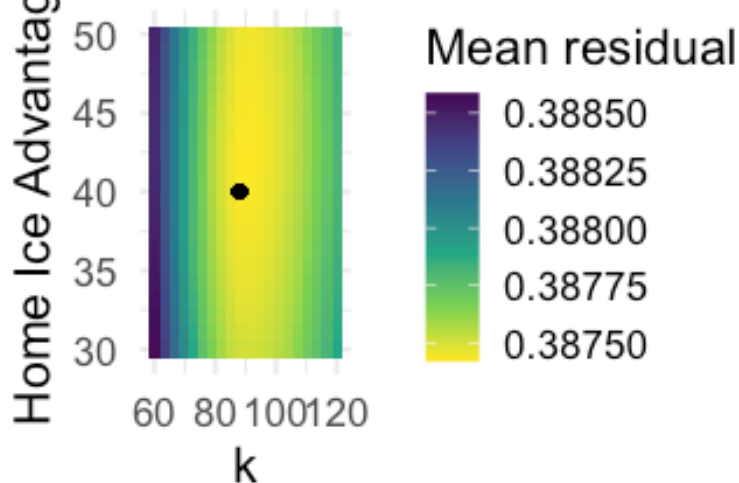
```
153 library(furrr)  
154 library(progressr)  
155 library(scico)  
156 plan(multisession)  
157 handlers("progress")  
158 options(progressr.enable = TRUE)  
159  
160 grid_100 = expand.grid(k = seq(60, 120, length.out = 20), home_ice = seq(30,  
161 50, length.out = 20), d = seq(0, 100, length.out = 1))  
162  
163 mean_residuals_100 = with_progress({future_pmap_dbl(grid_100, \ (k, home_ice,  
164 d) update_rankings_residuals(season = schedule_reg, end_date = "2025-03-25",  
165 ratings = rankings2324, k = k, home_ice = home_ice, d = d), .progress =  
166 TRUE)})  
167  
168 residual_100_df <- grid_100 |> mutate(mean_residual = mean_residuals_100)  
169  
170 ggplot(data = residual_100_df, aes(x = k,  
171 y = home_ice)) +  
172   geom_tile(aes(fill = mean_residual)) +  
173   geom_point(aes(x = 88, y = 40), color = "black", fill = "black") +  
174   scale_fill_viridis_c(option = "D",  
175                         ##limits = c(0.36, 0.41),  
176                         oob = scales::squish,  
177                         name = "Mean residual",
```

```

178         direction = -1) +
179     labs(x = "k",
180          y = "Home Ice Advantage",
181          title = "Optimization of k and Home Ice Advantage",
182          caption = "k is optimized at 88, home ice advantage at 40, \nand goal
183 differential factor \nof 0.6686 * log(abs(score_diff)) + 0.8048") +
184     theme_minimal(base_size = 16) +
185     theme(legend.position = "right",
186           plot.title = element_text(hjust = 0.5),
187           plot.caption = element_text(hjust = 0.5),
188           plot.margin = margin(t = 40, b = 20, l = 60, r = 60))

```

Optimization of k and Home Ice Advantage



k is optimized at 88, home ice advantage at 40,
and goal differential factor
of $0.6686 \cdot \log(\text{abs}(\text{score_diff})) + 0.8048$

189

190 Heat map showing the grid search to optimize update factor, k, and home ice advantage,
 191 home_ice. k and home_ice are optimized to lower the mean absolute value of the
 192 difference between expected outcome and actual outcome. Goal differential factor was
 193 held at $(0.6686 \cdot \ln(\text{scoreDiff})) + 0.8048$. k = 88, home_ice = 40.

194 Plot of residuals:

```


195 apr15 = apr_15_ranking |> bind_rows()
196
197 schedule_elo = schedule |>
198   mutate(home_elo = NA) |>

```

```

199     mutate(away_elo = NA)
200
201 schedule_apr15 = left_join(schedule_elo, apr15,
202                             by = join_by(date == date, home_team == Team)) |>
203     mutate(home_elo = rating) |>
204     select(-rating)
205
206 merged_sched_apr15 = left_join(schedule_apr15, apr15,
207                                 by = join_by(date == date, away_team == Team)) |>
208     mutate(away_elo = rating) |>
209     select(-rating)
210
211 schedule_full_apr15 = merged_sched_apr15 |>
212     mutate(outcome_away = abs(outcome - 1)) |>
213     ## Calculating expected outcome variable for home and away team
214     mutate(exp_home = 1/(1 + 10^((away_elo - home_elo)/400))) |>
215     mutate(exp_away = 1/(1 + 10^((home_elo - away_elo)/400))) |>
216     ## Using expected outcome variable to generate new Elo ratings based on
217 actual outcome and expected outcome
218     mutate(elo_new_home = home_elo + 100 * (outcome - exp_home)) |>
219     mutate(elo_new_away = away_elo + 100 * (outcome_away - exp_away))
220
221 ## Making a residual column
222 schedule_full_apr15 <- schedule_full_apr15 |>
223     mutate(residual = outcome - exp_home) |>
224     mutate(abs_residual = abs(residual))
225
226 ggplot(data = schedule_full_apr15, aes(x = residual)) +
227     geom_histogram(color = "purple4", fill = "seagreen3") +
228     labs(caption = "Positive residual past 0.5 indicates model predicting a
229 home loss when actual result is a home win. Negative residual beyond -0.5
230 indicate a predicted home wins with an observed home loss") +
231     theme_minimal()

```



Negative residual beyond -0.5 indicate a predicted home wins with an observed home loss

Plot of residuals shows all the differences between expected outcome and actual outcome. Positive residual past 0.5 indicates model predicting a home loss when actual result is a home win. Negative residual beyond -0.5 indicate a predicted home wins with an observed home loss. We see that the model has large amounts of residuals around the $|0.25|$ which means that the expected outcome is still within the range of the actual outcome. For example, actual outcome is 0, and expected is 0.25, the model still predicted a loss. The only grey area is around ties, since ties are 0.5, a $|0.25|$ difference could be construed as tie or a win/loss.

Binned expected value graph:

```
schedule_full_apr15 |> summarise(mean_resid = mean(abs_residual, na.rm =
TRUE))
```

```
# A tibble: 1 × 1
  mean_resid
    <dbl>
1      0.294
```

```
prop_wins15 <- schedule_full_apr15 |>
  mutate(binned_exp = floor(exp_home / 0.1) * 0.1 + 0.05) |>
  group_by(binned_exp) |>
  summarise(win_prop = mean(outcome, na.rm = TRUE),
```



```

252     totalgames = n()) |>
253     filter(!is.na(binned_exp))
254
255 modgdha = lm(win_prop ~ binned_exp, data = prop_wins15, weights = totalgames)
256 summary(modgdha)

```

Call:

```
lm(formula = win_prop ~ binned_exp, data = prop_wins15, weights = totalgames)
```

Weighted Residuals:

	Min	1Q	Median	3Q	Max
	-1.03772	-0.41495	0.01982	0.49843	0.79635

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.11495	0.04086	-2.813	0.0227 *
binned_exp	1.24411	0.06996	17.782	1.02e-07 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6344 on 8 degrees of freedom
Multiple R-squared: 0.9753, Adjusted R-squared: 0.9722
F-statistic: 316.2 on 1 and 8 DF, p-value: 1.024e-07

```

275 ggplot(data = prop_wins15, aes(x = binned_exp,
276                               y = win_prop,
277                               size = totalgames)) +
278   geom_point(color = "black", shape = 16) +
279   geom_smooth(aes(color = "Fitted Model",
280                  weight = totalgames), method = "lm", se = FALSE, size =
281 1.2) +
282   geom_abline(data = data.frame(1),
283              aes(color = "Expected Linear Model",
284                  linetype = "Expected Linear Model"),
285              slope = 1, intercept = 0, linetype = 2, size = 1) +
286   scale_color_manual(name = "Model",
287                     values = c("Fitted Model" = "gold",
288                                "Expected Linear Model" = "black")) +
289   labs(title = "Proportion of Home Team Wins \nfrom Home Expected Outcome",
290        x = "Expected Outcome",
291        y = "Proportion of \nGames Won",
292        caption = "Size of points indicate more games played. \nModel weighs
293 point based off of amount of games played") +
294   guides(size = "none") +
295   theme_minimal(base_size = 16) +
296   theme(legend.position = "right",
297         legend.background = element_rect(fill = "white", color = NA),
298         legend.title = element_text(size = 12),
299         legend.text = element_text(size = 10),

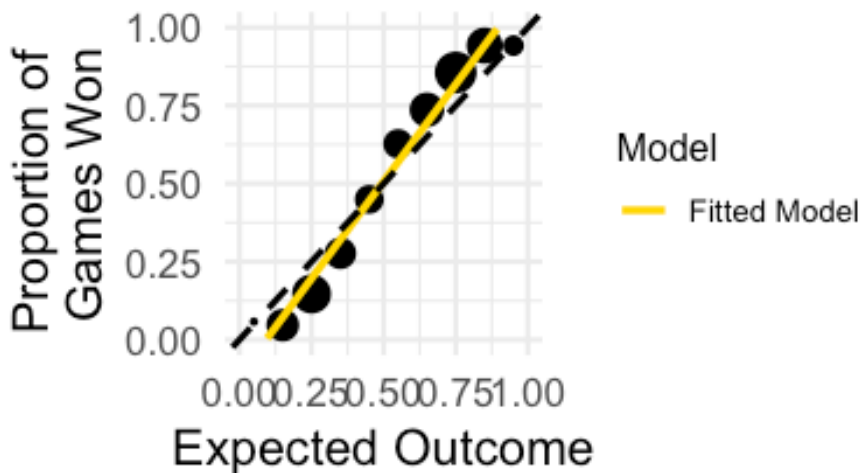
```

```

300     plot.title = element_text(hjust = 0.5),
301     plot.caption = element_text(hjust = 0.5),
302     plot.margin = margin(t = 40, b = 20, l = 40, r = 40)) +
303     xlim(c(0, 1)) +
304     ylim(c(0, 1))

```

Proportion of Home Team Wins from Home Expected Outcome



Size of points indicate more games played.
Model weighs point based off of amount of games played

305
 306 This plot bins every game played in the 2024-2025 season by every 0.1. Then uses the
 307 amount of games played in each bin to weigh each point. The y-axis is the proportion of
 308 wins for each binned proportion. We expect a slope of 1. This means that for example, at
 309 expected outcome of 0.9 we expect 90% of the games are won. In this case the slope of our
 310 fitted model is 1.24 with an intercept of -0.11. This further shows the accuracy of the
 311 optimized Elo function.

312 Final rankings are shown below:

```

313 apr15 = apr_15_ranking |> bind_rows()
314
315 apr15_lagged = apr15 |> group_by(date) |>
316   summarise(last_date = last(date)) |>
317   mutate(lag_date = lag(last_date)) |>
318   select(-last_date)
319
320 apr15_full = left_join(apr15, apr15_lagged, join_by(date == lag_date)) |>


```

```

321 select(-date) |>
322 rename(date = date.y)
323
324 highlight = c("St. Lawrence", "Western Michigan", "Clarkson", "Bentley",
325 "RIT", "Cornell")
326
327 highlighted_color = c(
328   "St. Lawrence" = "firebrick",
329   "Western Michigan" = "goldenrod",
330   "Clarkson" = "forestgreen",
331   "Bentley" = "royalblue",
332   "RIT" = "darkorange",
333   "Cornell" = "red")
334
335 apr15_color = apr15_full |> mutate(highlight = if_else(Team %in% highlight,
336 Team, "Other"))
337
338 ggplot(data = apr15_color, aes(x = date,
339                                y = rating,
340                                group = Team)) +
341   geom_line(aes(color = highlight,
342                 alpha = highlight,
343                 linewidth = highlight)) +
344   scale_color_manual(values = c("Other" = "grey80",
345                                 highlighted_color),
346                      name = "Team") +
347   scale_alpha_manual(values = c("Other" = 0.5,
348                                 setNames(rep(1, length(highlighted_color)),
349                                           names(highlighted_color))),
350                      guide = "none") +
351   scale_linewidth_manual(values = c("Other" = 0.2,
352                                     setNames(rep(1,
353 length(highlighted_color)),
354                                               names(highlighted_color))),
355                          guide = "none") +
356   theme_minimal() +
357   labs(color = "Team",
358        title = "Full Season Rankings",
359        x = "Date",
360        y = "Rating") +
361   theme(legend.position = "right",
362         plot.title = element_text(hjust = 0.5))

```

Full Season Rankings



363

```
364 print(n = 64, apr_15_ranking[[110]] |> arrange(desc(rating)))
```

365 # A tibble: 64 × 3

	Team	rating	date
	<chr>	<dbl>	<date>
368	1 Western Michigan	2186.	2025-04-12
369	2 Boston University	1990.	2025-04-12
370	3 Connecticut	1970.	2025-04-12
371	4 Penn State	1964.	2025-04-12
372	5 Denver	1954.	2025-04-12
373	6 Michigan State	1943.	2025-04-12
374	7 Boston College	1915.	2025-04-12
375	8 Cornell	1909.	2025-04-12
376	9 Minnesota State	1902.	2025-04-12
377	10 Maine	1890.	2025-04-12
378	11 Massachusetts	1888.	2025-04-12
379	12 North Dakota	1829.	2025-04-12
380	13 Minnesota	1798.	2025-04-12
381	14 Quinnipiac	1796.	2025-04-12
382	15 Arizona State	1786.	2025-04-12
383	16 Clarkson	1785.	2025-04-12
384	17 Omaha	1774.	2025-04-12
385	18 Bentley	1766.	2025-04-12

386	19 St. Thomas	1744.	2025-04-12
387	20 Providence	1734.	2025-04-12
388	21 Michigan	1733.	2025-04-12
389	22 Northeastern	1730.	2025-04-12
390	23 Dartmouth	1706.	2025-04-12
391	24 Notre Dame	1697.	2025-04-12
392	25 Ohio State	1687.	2025-04-12
393	26 Long Island	1662.	2025-04-12
394	27 New Hampshire	1648.	2025-04-12
395	28 Holy Cross	1638.	2025-04-12
396	29 Harvard	1632.	2025-04-12
397	30 Bowling Green	1593.	2025-04-12
398	31 Colorado College	1575.	2025-04-12
399	32 Bemidji State	1574.	2025-04-12
400	33 Colgate	1563.	2025-04-12
401	34 Minnesota-Duluth	1558.	2025-04-12
402	35 Wisconsin	1549.	2025-04-12
403	36 Alaska	1538.	2025-04-12
404	37 Vermont	1536.	2025-04-12
405	38 Augustana	1533.	2025-04-12
406	39 St. Cloud State	1533.	2025-04-12
407	40 Brown	1530.	2025-04-12
408	41 Sacred Heart	1518.	2025-04-12
409	42 Merrimack	1514.	2025-04-12
410	43 Mass.-Lowell	1511.	2025-04-12
411	44 Union	1500.	2025-04-12
412	45 Princeton	1479.	2025-04-12
413	46 Army	1477.	2025-04-12
414	47 Lake Superior	1449.	2025-04-12
415	48 Ferris State	1444.	2025-04-12
416	49 Air Force	1435.	2025-04-12
417	50 Michigan Tech	1424.	2025-04-12
418	51 Lindenwood	1420.	2025-04-12
419	52 American Int'l	1407.	2025-04-12
420	53 Stonehill	1372.	2025-04-12
421	54 Niagara	1369.	2025-04-12
422	55 Rensselaer	1359.	2025-04-12
423	56 Canisius	1352.	2025-04-12
424	57 Alaska-Anchorage	1345.	2025-04-12
425	58 RIT	1326.	2025-04-12
426	59 Northern Michigan	1298.	2025-04-12
427	60 Yale	1281.	2025-04-12
428	61 St. Lawrence	1271.	2025-04-12
429	62 Robert Morris	1183.	2025-04-12
430	63 Miami	1177.	2025-04-12
431	64 Mercyhurst	1070.	2025-04-12

432 1. In these final season ranking our Elo system ranked Western Michigan as the top
433 team, with Boston University second. In fact, of the top 4 teams in our model, 3 of

them made the Frozen Four, with Denver finishing in 5th. Western Michigan ultimately ended up winning the National Championship.

7 Conclusion

In order to make a better ratings system in NCAA Men's Division 1 Ice Hockey, an Elo style system was created. Using goal differential, an update factor, and home ice advantage, an Elo model was optimized using a grid search method to get the mean absolute value of game residuals down to 0.388, and successfully ranked the four teams to make the Frozen Four in the top five and successfully ranked the national champions, Western Michigan, in the top spot for the end of 2024-2025 season rankings.

This model allows teams to be ranked quantitatively and use factors such as strength of schedule and quality of win/loss to accurately scale the effect of each win or loss. In the future, I would like to try to add more parameters in to the model to see how big of an effect things like, OT and neutral site have on expected outcome. Mostly I would like to break down score differential into goals for and goals against to see if having good defense is more important than good offense and vice versa. My biggest regret is not being able to optimize the score differential parameter myself, and if given more time would do this.

8 Code Appendix

Libraries:

```
library(elo)
library(dplyr)
library(tidyverse)
library(cowplot)
library(ggrepel)
library(lubridate)
library(rvest)
library(here)
library(forcats)
library(progressr)
library(furrr)
library(vctrs)
library(purrr)
```

Scraping function:

```
##Function to load in schedule
scrape_men <- function(season = "20232024"){
  ## URL for schedule data frame
  url_hockey <- paste("https://www.collegehockeynews.com/schedules/?season=",
season, sep = "")
  ## Selecting which schedule table to grab
  tab_hockey <- read_html(url_hockey) |>
    html_nodes("table")
}
```

```

475     ## The website likes to switch which table it uses. If function doesn't
476     work try changing which table number you select
477     stats_dirty <- tab_hockey[[1]] |> html_table()
478
479     ## Creating regex for date, and conference to make date and conference
480     columns in dataframe
481     regex_date <- "October|November|December|January|February|March|April"
482     regex_conference <- "Atlantic Hockey|Big Ten|CCHA|ECAC|Hockey
483     East|NCHC|Ind|Exhibition|Non-Conference"
484     ## Combining regexs with original table so that the original scraped
485     dataframe has date and conferene as variables
486     stats_regex <- stats_dirty |> mutate(date = if_else(str_detect(X1,
487     regex_date),
488
489                                     true = X1, false =
490     NA_character_),
491     conference = if_else(str_detect(X1,
492
493                                     true = X1, false
494     = NA_character_))
495
496     ## Filling in respective dates and conferences
497     stats_filled <- stats_regex |> fill(date, .direction = "down") |>
498     ## Selecting date and congference so they show up as X1 and X2 in the
499     dataframe
500     fill(conference, .direction = "down") |> select(date, conference,
501     everything())
502     ##filtering out anywhere that a conference value is undetected (Game
503     category, not an actual game played)
504     stats_filled_cleaner <- stats_filled |> filter(!str_detect(X1, regex_date)
505     &
506
507                                     !str_detect(X1,
508     regex_conference))
509     print(head(stats_filled_cleaner))
510
511     ## Dataframe is now in a format that is able to be worked on. Now creating
512     specific variables that we want to look at
513     ## Selecting first 8 columns
514     schedule_new <- stats_filled_cleaner |> select(date, conference, X1, X2,
515     X3, X4, X5, X6) |>
516     ## Taking out first two rows (no data in them). Renaming columns to match
517     what their variable is.
518     slice(-1 , -2) |> rename(game_type = conference, away_team = X1,
519     away_score = X2, location_marker = X3, home_team = X4, home_score = X5,
520     overtime = X6) |>
521     ## Take out the day of the week in our date columns as we don't need to
522     know if a game was played on a Monday per-se.
523     separate(col = date, into = c("weekday", "dm", "y"),
524
525                                     sep = ", ") |>
526     unite("new_date", c(dm, y),

```



```

524     sep = " ") |>
525     select(-weekday) |>
526     ##making date column into a <date> variable
527     mutate(date = mdy(new_date)) |>
528     ## Taking out the <chr> date variable
529     select(-new_date) |>
530     select(date, everything()) |>
531     ## Filtering out where there is no away team since that means no game was
532 played
533     filter(away_team != "") |>
534     ## Filtering out exhibition games since we aren't looking at exhibition
535 games
536     filter(game_type != "Exhibition") |>
537     ## Turning scores from <chr> to <dbl> variables
538     mutate(away_score = as.double(away_score)) |>
539     mutate(home_score = as.double(home_score)) |>
540     ## creating a variable to indicate if a game was played at a neutral site
541     mutate(neutral_site = case_when(location_marker == "vs." ~ 1,
542                                     location_marker == "at" ~ 0)) |>
543     ## Making the neutral_site variable as <lgl>
544     mutate(neutral_site = as.logical(neutral_site)) |>
545     ## taking out location_marker
546     select(-location_marker) |>
547     ## Making a Logical overtime variable. Note we are not differentiating
548 between OT and 2OT
549     mutate(overtime = case_when(overtime == "" ~ 0,
550                                 overtime == "ot" ~ 1,
551                                 overtime == "2ot" ~ 1)) |>
552     mutate(overtime = as.logical(overtime)) |>
553     ##Filtering out NA "overtime" values as this indicates no game played,
554 since overtime will either be TRUE or FALSE
555     filter(!is.na(overtime))|>
556     ## Creating a score differential variable to indicate a win, Loss, or tie
557 for the home team. If we know the outcome for the home team, we know the
558 outcome for the away team.
559     mutate(score_diff = home_score - away_score) |>
560     ## making an outcome variable for home team so ties get input as 0.5,
561 wins get input as 1, and Loss get input as 0.
562     mutate(outcome =
563             case_when(score_diff == 0 ~ "0.5",
564                       score_diff > 0 ~ "1",
565                       score_diff < 0 ~ "0")) |>
566     ## turning score_diff from <chr> to <dbl>
567     mutate(outcome = as.double(outcome)) |>
568     ## Filtering out games where D1 team played against D3 teams as these are
569 exhibition as well
570     filter(game_type != "Non-Conference v. D3")
571
572     ## Tidy schedule is returned
573     return(schedule_new)

```



```

574 }
575
576 ##Load in Schedule
577 schedule <- scrape_men("20242025")
578
579 ## Load in my arbitrary initial elo ranking
580 X22Rankings <- read_csv(here("datasets_dataframes/22Rankings.csv"))

```

581 Function to do single day ratings:

```

582 ##Function to update rankings
583 ##rating is the variable, ratings is the df.
584 update_rankings <- function(season, game_date, ratings, k = 20){
585   ## Filters schedule to a specific date
586   elo_ratings_update <- season |> filter(date == game_date) |>
587     ## Joins the Elo ratings from our rating file to the schedule file. Puts
588 updated ratings in the schedule
589   left_join(ratings, by = join_by(away_team == Team)) |>
590   rename(away_elo = rating) |>
591   ## Updates ratings for home team in the schedule file
592   left_join(ratings, by = join_by(home_team == Team)) |>
593   rename(home_elo = rating) |>
594   ## Creating an away team outcome variable. Opposite of home team or same
595 if tie.
596   mutate(outcome_away = abs(outcome - 1)) |>
597   ## Calculating expected outcome variable for home and away team
598   mutate(exp_home = 1/(1 + 10^((away_elo - home_elo)/400))) |>
599   mutate(exp_away = 1/(1 + 10^((home_elo - away_elo)/400))) |>
600   ## Using expected outcome variable to generate new Elo ratings based on
601 actual outcome and expected outcome
602   mutate(elo_new_home = home_elo + k*(outcome - exp_home)) |>
603   mutate(elo_new_away = away_elo + k*(outcome_away - exp_away)) |>
604   rename(date_update_rankings = date)
605
606   ranked_home <- left_join(ratings, elo_ratings_update, by = join_by(Team ==
607 home_team)) |>
608   relocate(elo_new_home) |>
609   mutate(rating = if_else(!is.na(elo_new_home),
610                          true = elo_new_home,
611                          false = rating)) |>
612   select(Team, rating, date_update_rankings)
613
614   ratings_new <- left_join(ranked_home, elo_ratings_update, by = join_by(Team
615 == away_team)) |>
616   relocate(elo_new_away) |>
617   mutate(rating = if_else(!is.na(elo_new_away),
618                          true = elo_new_away,
619                          false = rating)) |>
620   select(Team, rating, date_update_rankings.x) |>
621   mutate(date_update_rankings.x = game_date) |>

```

```

622     rename(date = date_update_rankings.x)
623
624     return(ratings_new)
625 }

```

626 Creating a vector of dates to be iterated to automate ratings:

```

627 dates_vec <- unique(schedule$date)
628 ## defining what new_rankings is going to be
629 new_rankings = rankings
630 ## Creating a for loop with the update_rankings function to update rankings
631 up to a specified date
632 for (i in dates_vec) {
633     new_rankings <- update_rankings(season = schedule, game_date = i, ratings =
634 new_rankings, k = 100)
635
636 }

```

637 Function to iterate through dates to automate ratings

```

638 update_rankings_iter <- function(season, end_date, ratings, k){
639
640     new_rankings <- list()
641
642     season_cut <- season |>
643     ## Filter by a specified end date to deal with NA values (gmaes that have
644 yet to be played)
645     filter(date <= ymd(end_date))
646     ## Creating a vector for unique dates in a season
647     dates_vector <- unique(season_cut$date)
648     ## Defining our rankings within the function
649     new_rankings[[1]] <- ratings
650     ## Creating a for loop for the function to generate new ratings with the
651 updtae_rankings function
652     for (i in 1:length(dates_vector)) {
653         new_rankings[[i + 1]] <- update_rankings(season = season_cut, game_date =
654 dates_vector[i], ratings = new_rankings[[i]], k = k)
655     }
656     return(new_rankings)
657 }

```

658 Function to update single day ratings with all three parameters, k, score differential, home
659 ice advantage:

```

660 update_rankings_gd_ha <- function(season, game_date, ratings, k = 100,
661 home_ice = 50, d = 0.5){
662     ## Filters schedule to a specific date
663     elo_ratings_update <- season |> filter(date == game_date) |>
664     ## Joins the Elo ratings from our rating file to the schedule file. Puts
665 updated ratings in the schedule

```

```

666 left_join(ratings, by = join_by(away_team == Team)) |>
667 rename(away_elo = rating) |>
668 ## Updates ratings for home team in the schedule file
669 left_join(ratings, by = join_by(home_team == Team)) |>
670 rename(home_elo = rating) |>
671 ## Creating an away team outcome variable. Opposite of home team or same
672 if tie.
673 mutate(outcome_away = abs(outcome - 1)) |>
674 ## Calculating expected outcome variable for home and away team
675 mutate(exp_home = 1/(1 + 10^((away_elo - (home_elo + home_ice))/400))) |>
676 mutate(exp_away = 1/(1 + 10^(((home_elo + home_ice) - away_elo)/400))) |>
677 ## Using expected outcome variable to generate new Elo ratings based on
678 actual outcome and expected outcome
679 ## fivethirtyeights parameters
680 ## 0.6686 * Log(abs(score_diff)) + 0.8048
681 mutate(score_mult = if_else(score_diff == 0,
682                               true = 0.8048,
683                               false = 0.6686 * log(abs(score_diff)) +
684 0.8048)) |>
685 mutate(elo_new_home = home_elo + k * score_mult * (outcome - exp_home))
686 |>
687 mutate(elo_new_away = away_elo + k * score_mult * (outcome_away -
688 exp_away)) |>
689 rename(date_update_rankings = date)
690
691 ranked_home_gd_ha <- left_join(ratings, elo_ratings_update, by =
692 join_by(Team == home_team)) |>
693 relocate(elo_new_home) |>
694 mutate(rating = if_else(!is.na(elo_new_home),
695                           true = elo_new_home,
696                           false = rating)) |>
697 select(Team, rating, date_update_rankings)
698
699 ratings_new_gd_ha <- left_join(ranked_home_gd_ha, elo_ratings_update, by =
700 join_by(Team == away_team)) |>
701 relocate(elo_new_away) |>
702 mutate(rating = if_else(!is.na(elo_new_away),
703                           true = elo_new_away,
704                           false = rating)) |>
705 select(Team, rating, date_update_rankings.x) |>
706 mutate(date_update_rankings.x = game_date) |>
707 rename(date = date_update_rankings.x)
708
709 return(ratings_new_gd_ha)
710 }

```

Function to iterate through entire season automatically:

```
rankings = X22Rankings
```

```

714 ##Function to update rankings
715 ##rating is the variable, ratings is the df.
716 update_rankings_gd_ha <- function(season, game_date, ratings, k = 100,
717 home_ice = 50, d = 0.5){
718   ## Filters schedule to a specific date
719   elo_ratings_update <- season |> filter(date == game_date) |>
720   ## Joins the Elo ratings from our rating file to the schedule file. Puts
721   updated ratings in the schedule
722   left_join(ratings, by = join_by(away_team == Team)) |>
723   rename(away_elo = rating) |>
724   ## Updates ratings for home team in the schedule file
725   left_join(ratings, by = join_by(home_team == Team)) |>
726   rename(home_elo = rating) |>
727   ## Creating an away team outcome variable. Opposite of home team or same
728   if tie.
729   mutate(outcome_away = abs(outcome - 1)) |>
730   ## Calculating expected outcome variable for home and away team
731   mutate(exp_home = 1/(1 + 10^((away_elo - (home_elo + home_ice))/400))) |>
732   mutate(exp_away = 1/(1 + 10^(((home_elo + home_ice) - away_elo)/400))) |>
733   ## Using expected outcome variable to generate new Elo ratings based on
734   actual outcome and expected outcome
735   ## 0.6686 * Log(abs(score_diff)) + 0.8048
736   mutate(score_mult = if_else(score_diff == 0,
737                               true = 0.8048,
738                               false = 0.6686 * log(abs(score_diff)) +
739                               0.8048)) |>
740   mutate(elo_new_home = home_elo + k * score_mult * (outcome - exp_home))
741   |>
742   mutate(elo_new_away = away_elo + k * score_mult * (outcome_away -
743   exp_away)) |>
744   ##mutate(elo_new_home = home_elo + k*(outcome - exp_home) + d *
745   score_diff) |>
746   mutate(elo_new_home = if_else(elo_new_home < 100,
747                               true = 100,
748                               false = elo_new_home)) |>
749   ##mutate(elo_new_away = away_elo + k*(outcome_away - exp_away) + d * -1 *
750   (score_diff)) |>
751   mutate(elo_new_away = if_else(elo_new_away < 100,
752                               true = 100,
753                               false = elo_new_away)) |>
754   rename(date_update_rankings = date)
755
756   ## Find out which is the right date, select() and relocate(), **DO THIS
757   FIRST**: try renaming date in one of the df to make less confusing (at
758   start).
759   ranked_home_gd_ha <- left_join(ratings, elo_ratings_update, by =
760   join_by(Team == home_team)) |>
761   relocate(elo_new_home) |>
762   mutate(rating = if_else(!is.na(elo_new_home),
763                           true = elo_new_home,

```

```

764         false = rating)) |>
765     select(Team, rating, date_update_rankings)
766
767     ratings_new_gd_ha <- left_join(ranked_home_gd_ha, elo_ratings_update, by =
768 join_by(Team == away_team)) |>
769     relocate(elo_new_away) |>
770     mutate(rating = if_else(!is.na(elo_new_away),
771                             true = elo_new_away,
772                             false = rating)) |>
773     select(Team, rating, date_update_rankings.x) |>
774     mutate(date_update_rankings.x = game_date) |>
775     rename(date = date_update_rankings.x)
776
777     return(ratings_new_gd_ha)
778 }

```

779 Creating a vector of dates to use again to automate:

```

780 dates_vec <- unique(schedule$date)
781 ## defining what new_rankings is going to be
782 new_rankings = rankings
783 ## Creating a for loop with the update_rankings function to update rankings
784 up to a specified date
785 for (i in dates_vec) {
786     new_rankings <- update_rankings_gd_ha(season = schedule, game_date = i,
787 ratings = new_rankings, k = 100, home_ice = 50, d = 0.5)
788 }
789 }

```

790 Function to return a completed set of rankings:

791 Function to run ratings and return the mean residual

```

792 update_rankings_residuals = function(season, end_date, ratings, k, home_ice,
793 d){
794
795     new_rankings = update_rankings_iter_gd_ha(season = season, end_date =
796 end_date, ratings = ratings, k = k, home_ice = home_ice, d = d)
797
798     full_rankings = new_rankings |> bind_rows()
799
800     lagged_dates = full_rankings |> group_by(date) |>
801     summarise(last_date = last(date)) |>
802     mutate(lag_date = lag(last_date)) |>
803     select(-last_date)
804
805     lagged_rankings <- left_join(full_rankings, lagged_dates, join_by(date ==
806 lag_date)) |>
807     select(-date) |>
808     rename(date = date.y)

```

```

809
810   season = season |>
811     mutate(home_elo = NA) |>
812     mutate(away_elo = NA)
813
814   merged_season_home = left_join(season, lagged_rankings,
815                                   by = join_by(date == date, home_team ==
816 Team)) |>
817     mutate(home_elo = rating) |>
818     select(-rating)
819
820   merged_season = left_join(merged_season_home, lagged_rankings,
821                                   by = join_by(date == date, away_team == Team))
822 |>
823     mutate(away_elo = rating) |>
824     select(-rating)
825
826   full_season = merged_season |>
827     mutate(outcome_away = abs(outcome - 1)) |>
828     ## Calculating expected outcome variable for home and away team
829     mutate(exp_home = 1/(1 + 10^((away_elo - (home_elo + home_ice))/400))) |>
830     mutate(exp_away = 1/(1 + 10^(((home_elo + home_ice) - away_elo)/400))) |>
831     ## Using expected outcome variable to generate new Elo ratings based on
832 actual outcome and expected outcome
833     ##
834     mutate(score_mult = if_else(score_diff == 0,
835                                   true = 0.8048,
836                                   false = 0.6686 * log(abs(score_diff)) +
837 0.8048)) |>
838     mutate(elo_new_home = home_elo + k * score_mult * (outcome - exp_home))
839 |>
840     mutate(elo_new_away = away_elo + k * score_mult * (outcome_away -
841 exp_away)) |>
842     mutate(elo_new_home = if_else(elo_new_home < 100,
843                                   true = 100,
844                                   false = elo_new_home)) |>
845     mutate(elo_new_away = if_else(elo_new_away < 100,
846                                   true = 100,
847                                   false = elo_new_away)) |>
848     mutate(residual = outcome - exp_home) |>
849     mutate(abs_residual = abs(residual))
850
851   mean_residual = full_season |> summarise(avg = mean(abs_residual, na.rm =
852 TRUE))
853
854   return(pull(mean_residual))
855 }

```

Additional data frames:

857 Optimization process:

858 Optimized ratings:

```
859 apr_15_ranking = update_rankings_iter_gd_ha(schedule, "2025-04-15",
860                                             rankings2324, optimal$k,
861                                             optimal$home_ice,
862                                             0)
```

863 Data frame for mapping plot:

```
864 plan(multisession)
865 handlers("progress")
866 options(progressr.enable = TRUE)
867
868 grid_100 = expand_grid(k = seq(60, 120, length.out = 20), home_ice = seq(30,
869 50, length.out = 20), d = seq(0, 100, length.out = 1))
870
871 mean_residuals_100 = with_progress({future_pmap_dbl(grid_100, \ (k, home_ice,
872 d) update_rankings_residuals(season = schedule_reg, end_date = "2025-03-25",
873 ratings = rankings2324, k = k, home_ice = home_ice, d = d), .progress =
874 TRUE)})
875
876 residual_100_df <- grid_100 |> mutate(mean_residual = mean_residuals_100)
```

877 9 Works Cited

878 College Hockey News. (2024). *Men's Division I Hockey Schedule – 2023–2024 Season*.
879 <https://www.collegehockeynews.com/schedules/>

880 College Hockey News. (2025). *Men's Division I Hockey Schedule – 2024–2025 Season*.
881 <https://www.collegehockeynews.com/schedules/>

882 **Silver, N.** (n.d.). *How our NHL predictions work*. FiveThirtyEight.
883 <https://fivethirtyeight.com/methodology/how-our-nhl-predictions-work/>

884

885 9.1