

1 Analyzing Elo Ratings For NCAA Men’s

2 Division 1 Hockey

3 Alex Tidd

4 Advisors: Matt Higham and Robin Lock

5 2025-05-05

6 Table of contents

7 1 Abstract..... 2

8 2 Introduction..... 2

9 3 Data 4

10 4 Elo Model 6

11 5 Results 9

12 6 Conclusion 21

13 7 Code Appendix 22

14 8 Works Cited 31

17 1 Abstract

18 There are 64 NCAA Division 1 Men's Hockey teams. In NCAA Division 1 Men's Hockey, the
19 U.S. College Hockey Online (USCHO) provides the official rankings on the 64 total teams
20 using a system that relies on expert votes. However, this method is not perfect as there is
21 no formal quantitative analytics involved in the voting. To improve on this, a chess ratings
22 system, called Elo, is commonly used in many sports to quantitatively rate players,
23 weighing strength of opponent and recency of match. In this project, we modify the Elo
24 system for NCAA Division 1 Men's Hockey by adding in weights for game goal differential
25 (so that games in which the score differential was large potentially result in a larger bump
26 in Elo for the winning team) and home-ice advantage (so that the team playing on home-ice
27 has an adjusted probability of winning the game).

28 2 Introduction

29 NCAA Men's Division 1 Ice Hockey has 64 teams. This inherently makes ranking teams
30 extremely hard. It is impossible to rank teams in the same manner as pro leagues like the
31 NHL and junior leagues like the 3 CHL leagues, where teams are ranked based off of a
32 record-point system. This is due to the fact that many teams end up with similar records,
33 but with a league so large and full of different conferences, the strength of a victory/loss is
34 not the same for every game. Some conferences are weaker than others, so if a weak-
35 conference playing team has a great record, it would not be fair to rank them higher than a
36 team in a strong conference with a mediocre record.

37 To rank teams, the current method is an "expert vote" system done by US College Hockey
38 Online (USCHO). This method takes votes by "experts" and ranks teams based on how
39 many votes each team receives. In general, these rankings are pretty accurate. However, it
40 lacks any true quantitative analysis, instead relying on the opinions of voters.

41 This project aims to solve the issue presented by using a ratings system used in chess to
42 accurately rate and rank players, called Elo, for use in collegiate hockey. In chess, Elo
43 ratings calculate an "expected outcome" using 2 players' ratings. This expected outcome

is essentially a “win probability” for each player. After the expected outcome is calculated, it is compared to the actual outcome. Depending on the outcome, a player’s new rating will either be increased (if the player won) or decreased (if the player lost). The benefit to using Elo, is that the rating system takes into consideration your strength, the opponent’s strength, and the recency of the match. Big wins/losses garner big adjustments in rating, whereas an expected win/loss won’t inflate/tarnish a rating very much. The equation for expected outcome is:

$$E_a = \frac{1}{1 + 10^{\frac{R_b - R_a}{400}}}$$

Where:

- E_a is the player’s expected outcome.
- R_a is the player’s pregame rating.
- R_b is the opponent’s pregame rating.

To update the rating after the match is played, the update equation is:

$$R'_a = R_a + k(\text{outcome} - E_a)$$

Where:

- R'_a is the player’s postgame rating.
- R_a is the player’s pregame rating
- k is an update factor that scales the amount of points a player changes by.
- outcome is the outcome of the match, 1 for a win, 0 for a loss.
- E_a is the player’s expected outcome.

Example:

65 Player A has a pregame rating of 2000 and beats Player B with a rating of 2200. To find
66 Player A's new rating (k is set to 100 for this example):

$$67 \quad E_a = \frac{1}{1 + 10^{\frac{2200 - 2000}{400}}}$$

$$68 \quad E_a = 0.2402531$$

$$69 \quad R'_a = 2000 + 100(1 - 0.2402531)$$

$$70 \quad R'_a = 2075.975$$

71 We see that Player A's new rating is 2075.975. It grew by almost 76 points due to them
72 beating a player they were not expecting to beat. Had Player A had a pregame rating off
73 2200 and Player B had a pregame rating of 2000, with Player A still winning, Player A's new
74 rating would be 2224.025, only a 24.025 point increase, as Player A was expected to win.


75 The goal for this project is to use this base Elo ratings system, incorporate factors such as
76 home ice advantage and goal differential, optimize k, and make a ratings model that can
77 accurately predict outcome and rank teams accordingly. This paper shows the steps of
78 finding data, creating a ratings function, and optimizing our constants.

79 3 Data

80 There were two main data sets used for this project. The full 2023-2024 NCAA Men's
81 Division 1 Ice Hockey schedule, which after wrangling, included 1166 games with
82 variables: date, game_type, away_team, away_score, home_team, home_score, overtime,
83 neutral_site, score_diff, and outcome. The score_diff and outcome variables are in
84 reference to the home team, with outcome being either 1: win, 0.5: tie, or 0: loss. The next
85 data frame was the entire 2024-2025 NCAA Men's Division 1 Ice Hockey schedule. The
86 schedule contains 1153 games which had the same variables and setup as the 2023-2024
87 season. Both data frames were scraped from the website: College Hockey News.

88 2023-2024 Season:


```

89 schedule2324 |> slice(1:4)
90 # A tibble: 4 × 10
91   date      game_type    away_team    away_score home_team home_score overt
92 ime
93   <date>    <chr>        <chr>        <dbl> <chr>        <dbl> <lgl>
94 1 2023-10-07 Non-Conference Lake Super...      2 Michigan...      5 FALSE
95 2 2023-10-07 Non-Conference Mass.-Lowe...      2 Alaska-A...      3 TRUE
96 3 2023-10-07 Non-Conference Denver          7 Alaska        3 FALSE
97 4 2023-10-07 Non-Conference Lindenwood        1 Air Force      4 FALSE
98 #  3 more variables: neutral_site <lgl>, score_diff <dbl>, outcome <dbl>

```

99 2024-2025 Season:

```

100 schedule |> slice(1:4)
101 # A tibble: 4 × 10
102   date      game_type    away_team    away_score home_team home_score overt
103 ime
104   <date>    <chr>        <chr>        <dbl> <chr>        <dbl> <lgl>
105 1 2024-10-04 Non-Conference Michigan S...      2 Lake Sup...      1 TRUE
106 2 2024-10-04 Non-Conference Minnesota ...      5 Michigan      2 FALSE
107 3 2024-10-04 Non-Conference Arizona St...      8 Air Force      1 FALSE
108 4 2024-10-04 Non-Conference Bowling Gr...      2 Mercyhur...      1 FALSE
109 #  3 more variables: neutral_site <lgl>, score_diff <dbl>, outcome <dbl>

```

110 A third data frame for initial rankings was created from scratch using knowledge from the
 111 final 2022-2023 season. Initial rankings were created by ranking teams from 2000 to 1300.
 112 The top 8 teams from the 2022-2023 season were assigned the rating of 2000, every 8
 113 teams, the assigned rating would decrease by 100. This initial ratings file was used as the
 114 initial ratings file to run with the 2023-2024 season to get starting ratings for the 2024-2025
 115 season.

116 To get these “accurate” ratings for the 2024-2025 schedule, the 2023-2024 schedule was
 117 run through the Elo function with $k = 100$, a value used based off a general exploration
 118 which gave relatively accurate rankings compared to the final 2023-2024 season. After, the
 119 ratings were scaled using the equation below:

$$120 \quad ratings_{pre} = (ratings_{final} \cdot 0.7) + 450$$

121 Where:

- $ratings_{pre}$ is the preseason ratings.

- $ratings_{final}$ is the final ratings from the previous season.

This adjustment of rankings is used to adjust for player turnover, new coaching hires, and off season improvement. The formula used is based off of the formula used by Fivethirtyeight in their NHL Elo ratings. These final adjusted ratings are what was used for the intial rating in the 2024-2025 season.

4 Elo Model

To add goal differential into the Elo system, it was determined that it should be placed in the “rating update” portion of the Elo system, due to goal differential being a “post-game” statistic. On the other hand, home ice advantage affects the predicted probability of a home team winning. Therefore, it was added to the “expected outcome calculation” portion of the Elo system. Home ice advantage was also incorporated as a simple point “boost” for the home team instead of a multiplier, based off of FiveThirtyEight’s NHL Elo model.

In order to obtain the “best” values for each parameter, a grid search method was used to optimize each of the three parameters to lower the mean absolute residual of the season. In this case, the “absolute residual” is calculated by the absolute value of the difference of the game outcome and the expected outcome of the home team. Essentially we want the smallest difference between actual outcome and expected outcome. The smaller the difference, the more accurate our function is predicting actual outcomes of games.

The un-optimized Elo function is as follows:

$$E_{home} = \frac{1}{1 + 10^{\frac{R_{away} - (R_{home} + homeIce)}{400}}}$$

$$E_{away} = \frac{1}{1 + 10^{\frac{(R_{home} + homeIce) - R_{away}}{400}}}$$

145 Where:

- 146 • E_{home} is the expected outcome of the home team.
- 147 • E_{away} is the expected outcome of the away team.
- 148 • R_{home} is the pregame rating of the home team.
- 149 • R_{away} is the pregame rating of the away team.
- 150 • $homeIce$ is the home ice advantage factor.

$$151 \quad R'_{home} = R_{home} + k(d(\text{scoreDiff}) + (\text{outcome} - E_{home}))$$

$$152 \quad R'_{away} = R_{away} + k(d(-\text{scoreDiff}) + (\text{outcome} - E_{away}))$$

153 Where:

- 154 • R'_{home} is the post game rating of the home team.
- 155 • R'_{away} is the post game rating of the away team.
- 156 • R_{home} is the pregame rating of the home team.
- 157 • R_{away} is the pregame rating of the away team.
- 158 • k is the update factor, scaling the amount of points a team gains from a win or loss.
- 159 • d is the goal differential factor.
- 160 • scoreDiff is the score differential in relation to the home team.
- 161 • outcome is the outcome of a game, 1 for a win, 0.5 for a tie, 0, for a loss.
- 162 • E_{home} is the expected outcome of the home team.
- 163 • E_{away} is the expected outcome of the away team.

164 Using the above model a general exploration into possible values of k , d , and homelce
165 could render a lower mean absolute residual. Based off the exploration, k is found at
166 roughly 100, d at roughly 50, and homelce at roughly 50. It should be noted that k was
167 explored first with $d = 0$, and $\text{homelce} = 0$. Next was homelce , which was found at varying
168 weights of k and $d = 0$. d was then explored with $k = 100$, and $\text{homelce} = 0$. The only
169 reason k is held in all three cases, is that if $k = 0$, the function would not be able to update
170 rating as the update portion would always be $R_{\text{home}} + 0$ or $R_{\text{away}} + 0$.

171 These explorations provided essential knowledge for the grid search as initial optimization
172 could be found using smaller ranges of values, significantly decreasing the run time of
173 each optimization. This background knowledge cut run time of one optimization down from
174 16.3 days to roughly 30 minutes as each optimization grid search used 1000 combinations
175 of the three variables instead of 4.5 million.

176 Using this method, k was found to be optimized at 37.22222, homelce optimized at
177 53.33333, and d optimized at 40.55556. This provided a mean absolute residual of
178 0.3515844. This looked very promising at first glance. However, upon deeper inspection,
179 two major concerns were found. First, the bottom four teams had end of season ratings in
180 the negatives. Second, there were many residuals of 1, -1, and 0. This raised the most
181 concern as expected outcome is bounded by $0 < \text{expected outcome} < 1$, meaning we
182 should never see residuals of exactly 1 or 0, unless the predicted outcome is exactly 1 or 0.
183 Upon deeper inspection, it was found that there were many games where expected
184 outcome was to the effect of 0.9 repeating or $1e-20$. While technically within our bounds,
185 are highly unrealistic due to the incredibly high and low probabilities. A good model that
186 predicts win probability for college hockey should never give a pregame win probability of
187 nearly 100% or 0%. Something needed to be changed with the model. Goal differential was
188 looked at first.

189 Upon inspecting and researching different ways to incorporate goal differential in an Elo
190 function, FiveThirtyEight was looked at again. In their NHL model, goal differential, d , was
191 incorporated as shown below:

$$(0.6686 \cdot \ln(\text{scoreDiff})) + 0.8048)$$

This approach is much more stable for a few reasons. First, a logarithmic function levels out at the upper and lower bounds. This mimics real life as there is much of a difference in blowout games. A 7-0 game and a 12-0 game are both considered relatively equal games as they're both blowouts. This stops teams from gaining huge amounts of points by really blowing out a team well after it is already "blown out". The function in addition to downplaying the effects of extremes, also contains all the ratings in the positive ratings. This stabilizes the ratings so we don't see extreme jumps in rating update from an upset win, and won't tank a team for being upset.

The new update function now looks as follows:

$$R'_{home} = R_{home} + k \left(\left((0.6686 \cdot \ln(\text{scoreDiff})) + 0.8048 \right) (\text{outcome} - E_{home}) \right)$$

$$R'_{away} = R_{away} + k \left(\left((0.6686 \cdot \ln(\text{scoreDiff})) + 0.8048 \right) (\text{outcome} - E_{away}) \right)$$

0.8048 is added into the d factor to account for games where $\text{scoreDiff} = 1$ since $\ln(1) = 0$ and $\ln(-1)$ is undefined. This means that any game where $\text{scoreDiff} = |1|$ d is equal to 0.8048.


Now that goal factor is optimized, all that was left was to rerun the grid search to optimize k and homeIc . Using the grid search method, k is optimized at 88 and homeIc optimized at 40, which gives a mean absolute residual of 0.3876531. Whilst this may not have a mean absolute residual lower than the first model, all teams maintain ratings above zero and keep realistic predicted probabilities.

5 Results

Results from this optimization show that $k = 88$, $\text{homeIc} = 40$, and the calculation to utilize $d = (0.6686 \cdot \ln(\text{scoreDiff})) + 0.8048$.

215 Below is the grid search for k and homeIce. Multiple grid searches were conducted based
216 on the explorations until the sequence was narrowed to a range of 10. The plot below
217 shows a large range for both k and d to aid in mapping visual.

```
218 library(furrr)
219 library(progressr)
220 library(scico)
221 plan(multisession)
222 handlers("progress")
223 options(progressr.enable = TRUE)
224
225 grid_100 = expand.grid(k = seq(60, 120, length.out = 20), home_ice = seq(30,
226 50, length.out = 20), d = seq(0, 100, length.out = 1))
227
228 mean_residuals_100 = with_progress({future_pmap_dbl(grid_100, \ (k, home_ice,
229 d) update_rankings_residuals(season = schedule_reg, end_date = "2025-03-25",
230 ratings = rankings2324, k = k, home_ice = home_ice, d = d), .progress = TRUE)
231 })
232
233 residual_100_df <- grid_100 |> mutate(mean_residual = mean_residuals_100)
234
235 ggplot(data = residual_100_df, aes(x = k,
236 y = home_ice)) +
237   geom_tile(aes(fill = mean_residual)) +
238   geom_point(aes(x = 88, y = 40), color = "black", fill = "black") +
239   scale_fill_viridis_c(option = "D",
240 oob = scales::squish,
241 name = "Mean residual",
242 direction = -1) +
243   labs(x = "k",
244 y = "Home Ice Advantage",
245 title = "Optimization of k and Home Ice Advantage") +
246   theme_minimal() +
247   theme(legend.position = "right",
248 plot.title = element_text(hjust = 0.5))
```




249
 250 The heat map showing the grid search to d , k , and homeIce . The factors are all optimized
 251 to lower the mean absolute value of the difference between expected outcome and actual
 252 outcome. In this case, the optimized values were able to lower the mean absolute value of
 253 the residuals to 0.3874257. d was held at $(0.6686 \cdot \ln(\text{scoreDiff})) + 0.8048$. $k = 88$,
 254 $\text{homeIce} = 40$.

255 After finding the optimized values for k and homeIce using $d = (0.6686 \cdot \ln(\text{scoreDiff})) +$
 256 0.8048 , the entire season was run using the `update_rankings_residuals()` function. This
 257 function allowed us to look at residuals for every game in the 2024-2025 season. To see
 258 how the function was behaving with respects to home teams and away teams, a plot of the
 259 residuals was looked at instead of mean residuals, which take away the ability to
 260 investigate how the function deals with home and away teams. A plot of the residuals of
 261 the 2024-2025 season is shown below.

```

262 apr15 = apr_15_ranking |> bind_rows()
263
264 schedule_elo = schedule |>
265   mutate(home_elo = NA) |>
266   mutate(away_elo = NA)
267
268 schedule_apr15 = left_join(schedule_elo, apr15,
269                             by = join_by(date == date, home_team == Team)) |>
270   mutate(home_elo = rating) |>
271   select(-rating)
272
273 merged_sched_apr15 = left_join(schedule_apr15, apr15,
274                                 by = join_by(date == date, away_team == Team)) |>
275   mutate(away_elo = rating) |>
276   select(-rating)
277
278 schedule_full_apr15 = merged_sched_apr15 |>
279   mutate(outcome_away = abs(outcome - 1)) |>
280   ## Calculating expected outcome variable for home and away team
281   mutate(exp_home = 1/(1 + 10^((away_elo - home_elo)/400))) |>
282   mutate(exp_away = 1/(1 + 10^((home_elo - away_elo)/400))) |>
283   ## Using expected outcome variable to generate new Elo ratings based on actual outcome and expected outcome
284   mutate(elo_new_home = home_elo + 100 * (outcome - exp_home)) |>
285   mutate(elo_new_away = away_elo + 100 * (outcome_away - exp_away))
286
287
288 ## Making a residual column
289 schedule_full_apr15 <- schedule_full_apr15 |>
290   mutate(residual = outcome - exp_home) |>
291   mutate(abs_residual = abs(residual))
292
293 ggplot(data = schedule_full_apr15, aes(x = residual)) +
294   geom_histogram(color = "purple4", fill = "seagreen3") +
295   theme_minimal()

```



296

297 The plot of residuals shows all the differences between expected outcome and actual
 298 outcome. The positive residuals past 0.5 indicates the the model is predicting a home loss
 299 when actual result is a home win. Negative residuals beyond -0.5 indicates that the model
 300 predicted a home win with an observed home loss. We see that the model has large
 301 amounts of residuals around the $|0.25|$ which means that the expected outcome is still
 302 within the range of the actual outcome. For example, actual outcome is 0, and expected is
 303 0.25, the model still predicted a loss. The only grey area is around ties, since ties are 0.5, a
 304 $|0.25|$ difference could be construed as tie or a win/loss.

305 After looking into the residuals, we wanted to test just how accurate the predicted win
 306 probabilities were. To do this, we took the entire season and binned each game by
 307 expected value in increments of 0.1. We then plotted the expected values against the
 308 proportion of games within the binned expectation that resulted in a home win. We expect
 309 a linear relation: $\hat{p}_{win} = \widehat{expected}_{binned}$. The expected relation is showing that of all teams

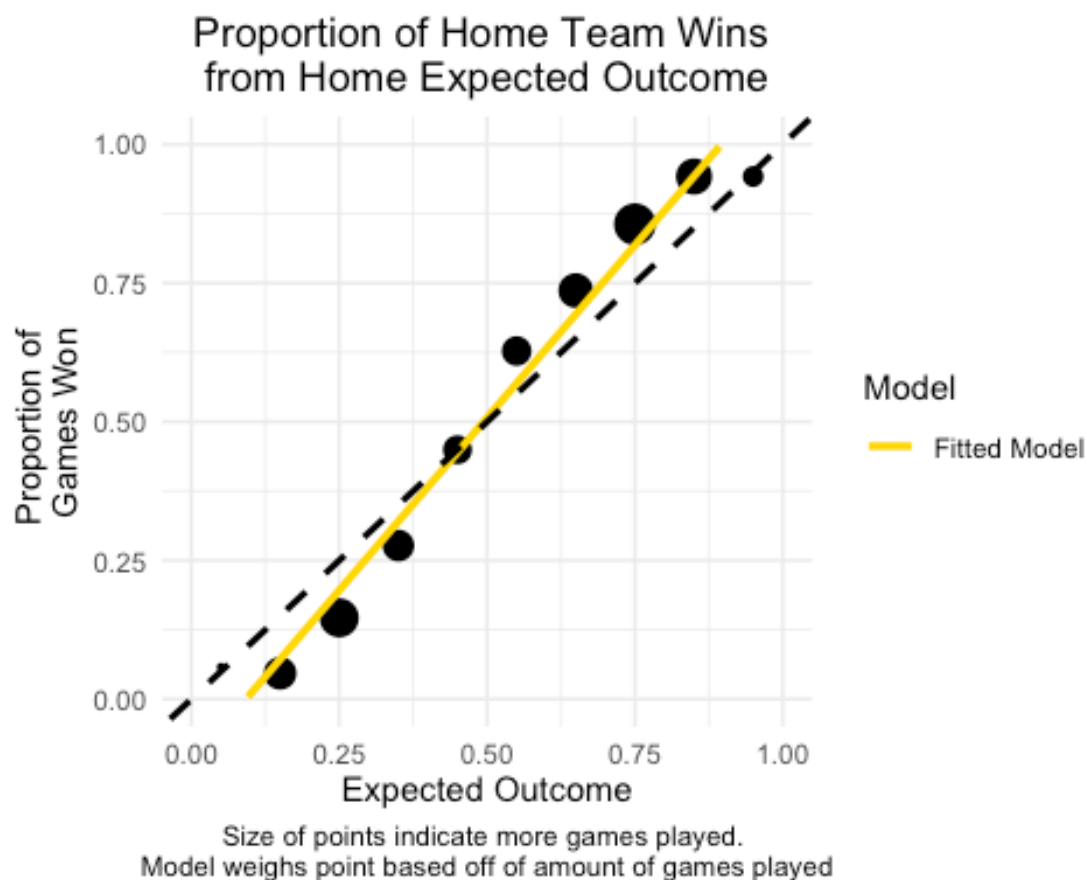
310 with a certain predicted win proportion, the proportion of times those teams win should be
311 equal to that predicted win proportion. Lastly, each binned expected value is weighted by
312 the amount of games in each bin so the linear model isn't skewed by bins with minimal
313 games. The plot is shown below.

```
314 schedule_full_apr15 |> summarise(mean_resid = mean(abs_residual, na.rm = TRUE  
315 ))  
316 # A tibble: 1 × 1  
317   mean_resid  
318     <dbl>  
319 1      0.294  
  
320 prop_wins15 <- schedule_full_apr15 |>  
321   mutate(binned_exp = floor(exp_home / 0.1) * 0.1 + 0.05) |>  
322   group_by(binned_exp) |>  
323   summarise(win_prop = mean(outcome, na.rm = TRUE),  
324             totalgames = n()) |>  
325   filter(!is.na(binned_exp))  
326  
327 modgdha = lm(win_prop ~ binned_exp, data = prop_wins15, weights = totalgames)  
328 summary(modgdha)  
  
329  
330 Call:  
331 lm(formula = win_prop ~ binned_exp, data = prop_wins15, weights = totalgames)  
332  
333 Weighted Residuals:  
334      Min       1Q   Median       3Q      Max  
335 -1.03772 -0.41495  0.01982  0.49843  0.79635  
336  
337 Coefficients:  
338             Estimate Std. Error t value Pr(>|t|)  
339 (Intercept) -0.11495    0.04086  -2.813   0.0227 *  
340 binned_exp   1.24411    0.06996  17.782 1.02e-07 ***  
341 ---  
342 Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
343  
344 Residual standard error: 0.6344 on 8 degrees of freedom  
345 Multiple R-squared:  0.9753,    Adjusted R-squared:  0.9722  
346 F-statistic: 316.2 on 1 and 8 DF,  p-value: 1.024e-07  
  
347 ggplot(data = prop_wins15, aes(x = binned_exp,  
348                               y = win_prop,  
349                               size = totalgames)) +  
350   geom_point(color = "black", shape = 16) +  
351   geom_smooth(aes(color = "Fitted Model",
```

```

352         weight = totalgames), method = "lm", se = FALSE, size = 1.2
353 ) +
354   geom_abline(data = data.frame(1),
355             aes(color = "Expected Linear Model",
356               linetype = "Expected Linear Model"),
357             slope = 1, intercept = 0, linetype = 2, size = 1) +
358   scale_color_manual(name = "Model",
359                     values = c("Fitted Model" = "gold",
360                               "Expected Linear Model" = "black")) +
361   labs(title = "Proportion of Home Team Wins \nfrom Home Expected Outcome",
362        x = "Expected Outcome",
363        y = "Proportion of \nGames Won",
364        caption = "Size of points indicate more games played. \nModel weighs p
365oint based off of amount of games played") +
366   guides(size = "none") +
367   theme_minimal() +
368   theme(legend.position = "right",
369         plot.title = element_text(hjust = 0.5),
370         plot.caption = element_text(hjust = 0.5)) +
371   xlim(c(0, 1)) +
372   ylim(c(0, 1))

```



This plot bins every game played in the 2024-2025 season by every 0.1. Then uses the amount of games played in each bin to weigh each point. The y-axis is the proportion of wins for each binned proportion. We expect a slope of 1. This means that for example, at expected outcome of 0.9 we expect 90% of the games are won. In this case the fitted model is $\hat{p}_{win} = 1.24411\widehat{expected}_{binned} - 1.24411$. This further shows the accuracy of the model.

Finally, we looked at the final rankings using the `update_rankings_iter_gd_ha()` function. This function runs the entire season and gathers ratings for each day and stores the rankings in a list. To see what the final rankings were, we pulled from the final day of the 2024-2025 post-season. To investigate the accuracy of the model, the rankings were compared to conference standings and the results of the NCAA Men's Division 1 Hockey National Championship Tournament. Final rankings are shown below with plot showing the progress of each team throughout the entire season.

```
apr15 = apr_15_ranking |> bind_rows()

apr15_lagged = apr15 |> group_by(date) |>
  summarise(last_date = last(date)) |>
  mutate(lag_date = lag(last_date)) |>
  select(-last_date)

apr15_full = left_join(apr15, apr15_lagged, join_by(date == lag_date)) |>
  select(-date) |>
  rename(date = date.y)

highlight = c("St. Lawrence", "Western Michigan", "Clarkson", "Bentley", "RIT",
  "Cornell")

highlighted_color = c(
  "St. Lawrence" = "firebrick",
  "Western Michigan" = "goldenrod",
  "Clarkson" = "forestgreen",
  "Bentley" = "royalblue",
  "RIT" = "darkorange",
  "Cornell" = "red")

apr15_color = apr15_full |> mutate(highlight = if_else(Team %in% highlight, T
eam, "Other"))
```




```

412 ggplot(data = apr15_color, aes(x = date,
413                               y = rating,
414                               group = Team)) +
415   geom_line(aes(color = highlight,
416                 alpha = highlight,
417                 linewidth = highlight)) +
418   scale_color_manual(values = c("Other" = "grey80",
419                                highlighted_color),
420                      name = "Team") +
421   scale_alpha_manual(values = c("Other" = 0.5,
422                                 setNames(rep(1, length(highlighted_color)),
423                                           names(highlighted_color))),
424                      guide = "none") +
425   scale_linewidth_manual(values = c("Other" = 0.2,
426                                     setNames(rep(1, length(highlighted_color)
427 ),
428                                               names(highlighted_color))),
429                          guide = "none") +
430   theme_minimal() +
431   labs(color = "Team",
432        title = "Full Season Rankings",
433        x = "Date",
434        y = "Rating") +
435   theme(legend.position = "right",
436         plot.title = element_text(hjust = 0.5))

```

Full Season Rankings



437

```
438 print(n = 64, apr_15_ranking[[110]] |> arrange(desc(rating))
439       |> select(-date))
```

440 # A tibble: 64 × 2

	Team	rating
	<chr>	<dbl>
443	1 Western Michigan	2186.
444	2 Boston University	1990.
445	3 Connecticut	1970.
446	4 Penn State	1964.
447	5 Denver	1954.
448	6 Michigan State	1943.
449	7 Boston College	1915.
450	8 Cornell	1909.
451	9 Minnesota State	1902.
452	10 Maine	1890.
453	11 Massachusetts	1888.
454	12 North Dakota	1829.
455	13 Minnesota	1798.
456	14 Quinnipiac	1796.
457	15 Arizona State	1786.
458	16 Clarkson	1785.
459	17 Omaha	1774.

460	18 Bentley	1766.
461	19 St. Thomas	1744.
462	20 Providence	1734.
463	21 Michigan	1733.
464	22 Northeastern	1730.
465	23 Dartmouth	1706.
466	24 Notre Dame	1697.
467	25 Ohio State	1687.
468	26 Long Island	1662.
469	27 New Hampshire	1648.
470	28 Holy Cross	1638.
471	29 Harvard	1632.
472	30 Bowling Green	1593.
473	31 Colorado College	1575.
474	32 Bemidji State	1574.
475	33 Colgate	1563.
476	34 Minnesota-Duluth	1558.
477	35 Wisconsin	1549.
478	36 Alaska	1538.
479	37 Vermont	1536.
480	38 Augustana	1533.
481	39 St. Cloud State	1533.
482	40 Brown	1530.
483	41 Sacred Heart	1518.
484	42 Merrimack	1514.
485	43 Mass.-Lowell	1511.
486	44 Union	1500.
487	45 Princeton	1479.
488	46 Army	1477.
489	47 Lake Superior	1449.
490	48 Ferris State	1444.
491	49 Air Force	1435.
492	50 Michigan Tech	1424.
493	51 Lindenwood	1420.
494	52 American Int'l	1407.
495	53 Stonehill	1372.
496	54 Niagara	1369.
497	55 Rensselaer	1359.
498	56 Canisius	1352.
499	57 Alaska-Anchorage	1345.
500	58 RIT	1326.
501	59 Northern Michigan	1298.
502	60 Yale	1281.
503	61 St. Lawrence	1271.
504	62 Robert Morris	1183.
505	63 Miami	1177.
506	64 Mercyhurst	1070.

507 In these final season rankings, our Elo system ranked Western Michigan as the top team,
508 with Boston University second. In fact, of the top 4 teams in our model, 3 of them made the
509 Frozen Four, with Denver finishing in 5th. Western Michigan ultimately ended up winning
510 the National Championship.

511 The progress chart shows some interesting patterns with some teams, especially Western
512 Michigan. They started the season middle of the pack, but with continuous strong wins,
513 they were able to reach 2nd place before the last quarter of the season. Other teams are
514 highlighted based on relation to St. Lawrence, such as Cornell (ECAC champions) and
515 Clarkson University (St. Lawrence's biggest rival). Other schools were chosen for personal
516 reasons, Bentley was chosen as they won their conference, Atlantic Hockey America for
517 the first time in the school's history and is also the school of my brother. RIT was also
518 chosen due to it being my hometown. One thing to note is how almost all of these schools
519 started in the middle of the rankings, but all finished in very different spots, showing just
520 how much a team's strength can change season to season. This further shows the
521 usefulness of adjusting team ratings prior to the season starting.

522 Finally, with the model optimized and set let's run an example, using the last game of the
523 year, the NCAA Championship game between Western Michigan and Boston University.

```
524 schedule_full_apr15 |> slice(1153) |>  
525   select(c(1, 3, 4, 5, 6, 9, 10, 11, 12))  
  
526 # A tibble: 1 × 9  
527   date       away_team      away_score home_team home_score score_diff outc  
528   <date>     <chr>          <dbl> <chr>         <dbl>     <dbl> <d  
529   1 2025-04-12 Western Michigan      6 Boston U...      2      -4  
530  
531 # 2 more variables: home_elo <dbl>, away_elo <dbl>
```

534 Western Michigan:

535

$$E_{WM} = \frac{1}{1 + 10^{\frac{(1990+40)-2186}{400}}}$$

536

$$E_{WM} = 0.71054$$

$$R'_{WM} = 2186 + 88 \left(\left((0.6686 \cdot \ln(4)) + 0.8048 \right) (1 - 0.71054) \right)$$

$$R'_{WM} = 2230.11$$

From this calculation, Western Michigan was a strong favorite to win this game, which they won handily, because of this their rating increased by 44 points to 2230.11. Now, let's see how Boston University was affected.

Boston University:

$$E_{BU} = \frac{1}{1 + 10^{\frac{2186 - (1990 + 40)}{400}}}$$

$$E_{BU} = 0.28946$$

$$R'_{BU} = 1990 + 88 \left(\left((0.6686 \cdot \ln(-4)) + 0.8048 \right) (1 - 0.28946) \right)$$

$$R'_{BU} = 1945.89$$

We see that again the change in rating is 44 with BU decreasing by 44 to 1945.89. This again, is a small decrease since BU were underdogs, and proceeded to lose by a large amount.

6 Conclusion

In order to make a better ratings system in NCAA Men's Division 1 Ice Hockey, an Elo style system was created. Using goal differential, an update factor, and home ice advantage, an Elo model was optimized using a grid search method to get the mean absolute value of game residuals down to 0.388. The model successfully ranked the four teams to make the Frozen Four in the top five and successfully ranked the national champions, Western Michigan, in the top spot for the end of 2024-2025 season rankings.

This model allows teams to be ranked quantitatively and use factors such as strength of schedule and quality of win/loss to accurately scale the effect of each win or loss. In the future, I would like to try to add more parameters in to the model to see how big of an effect

560 things like, OT and neutral site have on expected outcome. Mostly I would like to break
561 down score differential into goals for and goals against to see if having good defense is
562 more important than good offense and vice versa. My biggest regret is not being able to
563 optimize the score differential parameter myself, and if given more time, I would do this.

564 7 Code Appendix

565 Libraries:

```
566 library(elo)
567 library(dplyr)
568 library(tidyverse)
569 library(cowplot)
570 library(ggrepel)
571 library(lubridate)
572 library(rvest)
573 library(here)
574 library(forcats)
575 library(progressr)
576 library(furrr)
577 library(vctrs)
578 library(purrr)
```

579 Scraping function:

```
580 ##Function to load in schedule
581 scrape_men <- function(season = "20232024"){
582   ## URL for schedule data frame
583   url_hockey <- paste("https://www.collegehockeynews.com/schedules/?season=",
584     season, sep = "")
585   ## Selecting which schedule table to grab
586   tab_hockey <- read_html(url_hockey) |>
587     html_nodes("table")
588
589   ## The website likes to switch which table it uses. If function doesn't work
590   try changing which table number you select
591   stats_dirty <- tab_hockey[[1]] |> html_table()
592
593   ## Creating regex for date, and conference to make date and conference columns
594   in dataframe
595   regex_date <- "October|November|December|January|February|March|April"
596   regex_conference <- "Atlantic Hockey|Big Ten|CCHA|ECAC|Hockey East|NCHC|Ind
597 |Exhibition|Non-Conference"
598   ## Combining regexs with original table so that the original scraped dataframe
599   has date and conference as variables
```

```

600 stats_regex <- stats_dirty |> mutate(date = if_else(str_detect(X1, regex_da
601 te),
602                                     true = X1, false = NA_c
603 haracter_),
604                                     conference = if_else(str_detect(X1, re
605 gex_conference),
606                                     true = X1, false
607 = NA_character_))
608
609 ## Filling in respective dates and conferences
610 stats_filled <- stats_regex |> fill(date, .direction = "down") |>
611 ## Selecting date and conference so they show up as X1 and X2 in the dat
612 aframe
613 fill(conference, .direction = "down") |> select(date, conference, everyth
614 ing())
615 ##filtering out anywhere that a conference value is undetected (Game catego
616 ry, not an actual game played)
617 stats_filled_cleaner <- stats_filled |> filter(!str_detect(X1, regex_date)
618 &
619                                     !str_detect(X1, regex_conf
620 erence))
621 print(head(stats_filled_cleaner))
622
623 ## Dataframe is now in a format that is able to be worked on. Now creating
624 specific variables that we want to look at
625 ## Selecting first 8 columns
626 schedule_new <- stats_filled_cleaner |> select(date, conference, X1, X2, X3
627 , X4, X5, X6) |>
628 ## Taking out first two rows (no data in them). Renaming columns to match
629 what their variable is.
630 slice(-1, -2) |> rename(game_type = conference, away_team = X1, away_sco
631 re = X2, location_marker = X3, home_team = X4, home_score = X5, overtime = X6
632 ) |>
633 ## Take out the day of the week in our date columns as we don't need to k
634 now if a game was played on a Monday per-se.
635 separate(col = date, into = c("weekday", "dm", "y"),
636          sep = ", ") |>
637 unite("new_date", c(dm, y),
638       sep = " ") |>
639 select(-weekday) |>
640 ##making date column into a <date> variable
641 mutate(date = mdy(new_date)) |>
642 ## Taking out the <chr> date variable
643 select(-new_date) |>
644 select(date, everything()) |>
645 ## Filtering out where there is no away team since that means no game was
646 played
647 filter(away_team != "") |>
648 ## Filtering out exhibition games since we aren't looking at exhibition g
649 ames

```

```

650 filter(game_type != "Exhibition") |>
651   ## Turning scores from <chr> to <dbl> variables
652   mutate(away_score = as.double(away_score)) |>
653   mutate(home_score = as.double(home_score)) |>
654   ## creating a variable to indicate if a game was played at a neutral site
655   mutate(neutral_site = case_when(location_marker == "vs." ~ 1,
656                                   location_marker == "at" ~ 0)) |>
657   ## Making the neutral_site variable as <lgl>
658   mutate(neutral_site = as.logical(neutral_site)) |>
659   ## taking out location_marker
660   select(-location_marker) |>
661   ## Making a logical overtime variable. Note we are not differentiating be
662   tween OT and 2OT
663   mutate(overtime = case_when(overtime == "" ~ 0,
664                               overtime == "ot" ~ 1,
665                               overtime == "2ot" ~ 1)) |>
666   mutate(overtime = as.logical(overtime)) |>
667   ## Filtering out NA "overtime" values as this indicates no game played, si
668   nce overtime will either be TRUE or FALSE
669   filter(!is.na(overtime))|>
670   ## Creating a score differential variable to indicate a win, loss, or tie
671   for the home team. If we know the outcome for the home team, we know the outc
672   ome for the away team.
673   mutate(score_diff = home_score - away_score) |>
674   ## making an outcome variable for home team so ties get input as 0.5, win
675   s get input as 1, and loss get input as 0.
676   mutate(outcome =
677           case_when(score_diff == 0 ~ "0.5",
678                     score_diff > 0 ~ "1",
679                     score_diff < 0 ~ "0")) |>
680   ## turning score_diff from <chr> to <dbl>
681   mutate(outcome = as.double(outcome)) |>
682   ## Filtering out games where D1 team played against D3 teams as these are
683   exhibition as well
684   filter(game_type != "Non-Conference v. D3")
685
686   ## Tidy schedule is returned
687   return(schedule_new)
688 }
689
690 ##Load in Schedule
691 schedule <- scrape_men("20242025")
692
693 ## Load in my arbitrary initial elo ranking
694 X22Rankings <- read_csv(here("datasets_dataframes/22Rankings.csv"))

```

Function to do single day ratings:


```

696 ##Function to update rankings
697 ##rating is the variable, ratings is the df.
698 update_rankings <- function(season, game_date, ratings, k = 20){
699   ## Filters schedule to a specific date
700   elo_ratings_update <- season |> filter(date == game_date) |>
701     ## Joins the Elo ratings from our rating file to the schedule file. Puts
702     updated ratings in the schedule
703     left_join(ratings, by = join_by(away_team == Team)) |>
704     rename(away_elo = rating) |>
705     ## Updates ratings for home team in the schedule file
706     left_join(ratings, by = join_by(home_team == Team)) |>
707     rename(home_elo = rating) |>
708     ## Creating an away team outcome variable. Opposite of home team or same
709     if tie.
710     mutate(outcome_away = abs(outcome - 1)) |>
711     ## Calculating expected outcome variable for home and away team
712     mutate(exp_home = 1/(1 + 10^((away_elo - home_elo)/400))) |>
713     mutate(exp_away = 1/(1 + 10^((home_elo - away_elo)/400))) |>
714     ## Using expected outcome variable to generate new Elo ratings based on a
715     ctual outcome and expected outcome
716     mutate(elo_new_home = home_elo + k*(outcome - exp_home)) |>
717     mutate(elo_new_away = away_elo + k*(outcome_away - exp_away)) |>
718     rename(date_update_rankings = date)
719
720     ranked_home <- left_join(ratings, elo_ratings_update, by = join_by(Team ==
721     home_team)) |>
722     relocate(elo_new_home) |>
723     mutate(rating = if_else(!is.na(elo_new_home),
724     true = elo_new_home,
725     false = rating)) |>
726     select(Team, rating, date_update_rankings)
727
728     ratings_new <- left_join(ranked_home, elo_ratings_update, by = join_by(Team
729     == away_team)) |>
730     relocate(elo_new_away) |>
731     mutate(rating = if_else(!is.na(elo_new_away),
732     true = elo_new_away,
733     false = rating)) |>
734     select(Team, rating, date_update_rankings.x) |>
735     mutate(date_update_rankings.x = game_date) |>
736     rename(date = date_update_rankings.x)
737
738     return(ratings_new)
739 }

```

740 Creating a vector of dates to be iterated to automate ratings:

```

741 dates_vec <- unique(schedule$date)
742 ## defining what new_rankings is going to be

```

```

743 new_rankings = rankings
744 ## Creating a for Loop with the update_rankings function to update rankings u
745 p to a specified date
746 for (i in dates_vec) {
747   new_rankings <- update_rankings(season = schedule, game_date = i, ratings =
748 new_rankings, k = 100)
749 }
750

```

751 Function to iterate through dates to automate ratings

```

752 update_rankings_iter <- function(season, end_date, ratings, k){
753
754   new_rankings <- list()
755
756   season_cut <- season |>
757     ## Filter by a specified end date to deal with NA values (gmaes that have
758 yet to be played)
759   filter(date <= ymd(end_date))
760   ## Creating a vector for unique dates in a season
761   dates_vector <- unique(season_cut$date)
762   ## Defining our rankings within the function
763   new_rankings[[1]] <- ratings
764   ## Creating a for Loop for the function to generate new ratings with the up
765 dtae_rankings function
766   for (i in 1:length(dates_vector)) {
767     new_rankings[[i + 1]] <- update_rankings(season = season_cut, game_date =
768 dates_vector[i], ratings = new_rankings[[i]], k = k)
769   }
770   return(new_rankings)
771 }

```

772 Function to update single day ratings with all three parameters, k, score differential, home
773 ice advantage:

```

774 update_rankings_gd_ha <- function(season, game_date, ratings, k = 100, home_i
775 ce = 50, d = 0.5){
776   ## Filters schedule to a specific date
777   elo_ratings_update <- season |> filter(date == game_date) |>
778     ## Joins the Elo ratings from our rating file to the schedule file. Puts
779 updated ratings in the schedule
780   left_join(ratings, by = join_by(away_team == Team)) |>
781   rename(away_elo = rating) |>
782   ## Updates ratings for home team in the schedule file
783   left_join(ratings, by = join_by(home_team == Team)) |>
784   rename(home_elo = rating) |>
785   ## Creating an away team outcome variable. Opposite of home team or same
786 if tie.
787   mutate(outcome_away = abs(outcome - 1)) |>

```

```

788     ## Calculating expected outcome variable for home and away team
789     mutate(exp_home = 1/(1 + 10^((away_elo - (home_elo + home_ice))/400))) |>
790     mutate(exp_away = 1/(1 + 10^(((home_elo + home_ice) - away_elo)/400))) |>
791     ## Using expected outcome variable to generate new Elo ratings based on a
792     ctual outcome and expected outcome
793     ## fivethirtyrights parameters
794     ## 0.6686 * log(abs(score_diff)) + 0.8048
795     mutate(score_mult = if_else(score_diff == 0,
796                                true = 0.8048,
797                                false = 0.6686 * log(abs(score_diff)) + 0.804
798 8)) |>
799     mutate(elo_new_home = home_elo + k * score_mult * (outcome - exp_home)) |
800 >
801     mutate(elo_new_away = away_elo + k * score_mult * (outcome_away - exp_awa
802 y)) |>
803     rename(date_update_rankings = date)
804
805     ranked_home_gd_ha <- left_join(ratings, elo_ratings_update, by = join_by(Te
806 am == home_team)) |>
807     relocate(elo_new_home) |>
808     mutate(rating = if_else(!is.na(elo_new_home),
809                             true = elo_new_home,
810                             false = rating)) |>
811     select(Team, rating, date_update_rankings)
812
813     ratings_new_gd_ha <- left_join(ranked_home_gd_ha, elo_ratings_update, by =
814 join_by(Team == away_team)) |>
815     relocate(elo_new_away) |>
816     mutate(rating = if_else(!is.na(elo_new_away),
817                             true = elo_new_away,
818                             false = rating)) |>
819     select(Team, rating, date_update_rankings.x) |>
820     mutate(date_update_rankings.x = game_date) |>
821     rename(date = date_update_rankings.x)
822
823     return(ratings_new_gd_ha)
824 }

```

825 Function to iterate through entire season automatically:

```

826 rankings = X22Rankings
827
828 ##Function to update rankings
829 ##rating is the variable, ratings is the df.
830 update_rankings_gd_ha <- function(season, game_date, ratings, k = 100, home_i
831 ce = 50, d = 0.5){
832     ## Filters schedule to a specific date
833     elo_ratings_update <- season |> filter(date == game_date) |>
834     ## Joins the Elo ratings from our rating file to the schedule file. Puts

```

```

835 updated ratings in the schedule
836 left_join(ratings, by = join_by(away_team == Team)) |>
837 rename(away_elo = rating) |>
838 ## Updates ratings for home team in the schedule file
839 left_join(ratings, by = join_by(home_team == Team)) |>
840 rename(home_elo = rating) |>
841 ## Creating an away team outcome variable. Opposite of home team or same
842 if tie.
843 mutate(outcome_away = abs(outcome - 1)) |>
844 ## Calculating expected outcome variable for home and away team
845 mutate(exp_home = 1/(1 + 10^((away_elo - (home_elo + home_ice))/400))) |>
846 mutate(exp_away = 1/(1 + 10^(((home_elo + home_ice) - away_elo)/400))) |>
847 ## Using expected outcome variable to generate new Elo ratings based on a
848 ctual outcome and expected outcome
849 ## 0.6686 * Log(abs(score_diff)) + 0.8048
850 mutate(score_mult = if_else(score_diff == 0,
851                             true = 0.8048,
852                             false = 0.6686 * log(abs(score_diff)) + 0.804
853 8)) |>
854 mutate(elo_new_home = home_elo + k * score_mult * (outcome - exp_home)) |
855 >
856 mutate(elo_new_away = away_elo + k * score_mult * (outcome_away - exp_awa
857 y)) |>
858 ##mutate(elo_new_home = home_elo + k*(outcome - exp_home) + d * score_dif
859 f) |>
860 mutate(elo_new_home = if_else(elo_new_home < 100,
861                             true = 100,
862                             false = elo_new_home)) |>
863 ##mutate(elo_new_away = away_elo + k*(outcome_away - exp_away) + d * -1 *
864 (score_diff)) |>
865 mutate(elo_new_away = if_else(elo_new_away < 100,
866                             true = 100,
867                             false = elo_new_away)) |>
868 rename(date_update_rankings = date)
869
870 ## Find out which is the right date, select() and relocate(), **DO THIS FIR
871 ST**: try renaming date in one of the df to make less confusing (at start).
872 ranked_home_gd_ha <- left_join(ratings, elo_ratings_update, by = join_by(Te
873 am == home_team)) |>
874 relocate(elo_new_home) |>
875 mutate(rating = if_else(!is.na(elo_new_home),
876                         true = elo_new_home,
877                         false = rating)) |>
878 select(Team, rating, date_update_rankings)
879
880 ratings_new_gd_ha <- left_join(ranked_home_gd_ha, elo_ratings_update, by =
881 join_by(Team == away_team)) |>
882 relocate(elo_new_away) |>
883 mutate(rating = if_else(!is.na(elo_new_away),
884                         true = elo_new_away,

```

```

885         false = rating)) |>
886     select(Team, rating, date_update_rankings.x) |>
887     mutate(date_update_rankings.x = game_date) |>
888     rename(date = date_update_rankings.x)
889
890     return(ratings_new_gd_ha)
891 }

```

892 Creating a vector of dates to use again to automate:

```

893 dates_vec <- unique(schedule$date)
894 ## defining what new_rankings is going to be
895 new_rankings = ratings
896 ## Creating a for loop with the update_rankings function to update rankings u
897 p to a specified date
898 for (i in dates_vec) {
899     new_rankings <- update_rankings_gd_ha(season = schedule, game_date = i, rat
900 ings = new_rankings, k = 100, home_ice = 50, d = 0.5)
901 }
902 }

```

903 Function to return a completed set of rankings:

904 Function to run ratings and return the mean residual:

```

905 update_rankings_residuals = function(season, end_date, ratings, k, home_ice,
906 d){
907
908     new_rankings = update_rankings_iter_gd_ha(season = season, end_date = end_d
909 ate, ratings = ratings, k = k, home_ice = home_ice, d = d)
910
911     full_rankings = new_rankings |> bind_rows()
912
913     lagged_dates = full_rankings |> group_by(date) |>
914         summarise(last_date = last(date)) |>
915         mutate(lag_date = lag(last_date)) |>
916         select(-last_date)
917
918     lagged_rankings <- left_join(full_rankings, lagged_dates, join_by(date == l
919 ag_date)) |>
920     select(-date) |>
921     rename(date = date.y)
922
923     season = season |>
924     mutate(home_elo = NA) |>
925     mutate(away_elo = NA)
926
927     merged_season_home = left_join(season, lagged_rankings,

```

```

928         by = join_by(date == date, home_team == Team
929 )) |>
930   mutate(home_elo = rating) |>
931   select(-rating)
932
933   merged_season = left_join(merged_season_home, lagged_rankings,
934                             by = join_by(date == date, away_team == Team)) |
935   >
936   mutate(away_elo = rating) |>
937   select(-rating)
938
939   full_season = merged_season |>
940   mutate(outcome_away = abs(outcome - 1)) |>
941   ## Calculating expected outcome variable for home and away team
942   mutate(exp_home = 1/(1 + 10^((away_elo - (home_elo + home_ice))/400))) |>
943   mutate(exp_away = 1/(1 + 10^(((home_elo + home_ice) - away_elo)/400))) |>
944   ## Using expected outcome variable to generate new Elo ratings based on a
945 ctual outcome and expected outcome
946   ##
947   mutate(score_mult = if_else(score_diff == 0,
948                               true = 0.8048,
949                               false = 0.6686 * log(abs(score_diff)) + 0.804
950 8)) |>
951   mutate(elo_new_home = home_elo + k * score_mult * (outcome - exp_home)) |
952   >
953   mutate(elo_new_away = away_elo + k * score_mult * (outcome_away - exp_awa
954 y)) |>
955   mutate(elo_new_home = if_else(elo_new_home < 100,
956                                 true = 100,
957                                 false = elo_new_home)) |>
958   mutate(elo_new_away = if_else(elo_new_away < 100,
959                                 true = 100,
960                                 false = elo_new_away)) |>
961   mutate(residual = outcome - exp_home) |>
962   mutate(abs_residual = abs(residual))
963
964   mean_residual = full_season |> summarise(avg = mean(abs_residual, na.rm = T
965 RUE))
966
967   return(pull(mean_residual))
968 }

```

969 Additional data frames:

970 Optimization process:

971 Optimized ratings:

```
972 apr_15_ranking = update_rankings_iter_gd_ha(schedule, "2025-04-15",
973                                             rankings2324, optimal$k, optimal$
974 home_ice,
975                                             0)
```

976 Data frame for mapping plot:

```
977 plan(multisession)
978 handlers("progress")
979 options(progressr.enable = TRUE)
980
981 grid_100 = expand.grid(k = seq(60, 120, length.out = 20), home_ice = seq(30,
982 50, length.out = 20), d = seq(0, 100, length.out = 1))
983
984 mean_residuals_100 = with_progress({future_pmap_dbl(grid_100, \ (k, home_ice,
985 d) update_rankings_residuals(season = schedule_reg, end_date = "2025-03-25",
986 ratings = rankings2324, k = k, home_ice = home_ice, d = d), .progress = TRUE)
987 })
988
989 residual_100_df <- grid_100 |> mutate(mean_residual = mean_residuals_100)
```

990 8 Works Cited

991 College Hockey News. (2024). *Men's Division I Hockey Schedule – 2023–2024 Season*.
992 <https://www.collegehockeynews.com/schedules/>

993 College Hockey News. (2025). *Men's Division I Hockey Schedule – 2024–2025 Season*.
994 <https://www.collegehockeynews.com/schedules/>

995 **Silver, N.** (n.d.). *How our NHL predictions work*. FiveThirtyEight.
996 <https://fivethirtyeight.com/methodology/how-our-nhl-predictions-work/>

997
